

Reasoning Distillation: Qwen2.5-32b (base) vs DeepSeek-R1-32b (distilled) — with 95% Bootstrap CI

Model	MAE ↓	RMSE ↓	Correlation ↑	Kappa ↑	WD ↓	Stability ↓
Qwen2.5-32b (base) — Fewshot	0.66 [0.47, 0.87]	1.02 [0.78, 1.26]	0.46 [0.18, 0.67]	0.43 [0.15, 0.64]	0.40 [0.21, 0.64]	0.405
Qwen2.5-32b (base) — Zeroshot	0.53 [0.38, 0.66]	0.75 [0.63, 0.87]	0.55 [0.37, 0.69]	0.50 [0.33, 0.64]	0.34 [0.17, 0.51]	0.588
DeepSeek-R1-32b (distilled) — Fewshot	0.53 [0.36, 0.72]	0.85 [0.66, 1.02]	0.56 [0.32, 0.73]	0.54 [0.29, 0.71]	0.26 [0.17, 0.47]	0.449
DeepSeek-R1-32b (distilled) — Zeroshot	0.81 [0.60, 1.02]	1.09 [0.90, 1.26]	0.37 [0.13, 0.56]	0.30 [0.10, 0.47]	0.58 [0.36, 0.83]	0.658